

YIWEI LIU

🏡 Lausanne, Switzerland

✉️ yiw.liu@epfl.ch

🔗 ivyyyyw.github.io

✉️ [Yiwei Liu](https://yiwei.liu)

🔗 [Ivyyyyw](https://ivyyyyw)

EDUCATION

École Polytechnique Fédérale de Lausanne (EPFL)

MSc in Digital Humanities (GPA: 5.3/6)

Lausanne, Switzerland

Sept 2023 – Now

National Institute of Informatics

Visiting Student, advised by Prof. Saku Sugawara

Tokyo, Japan

Oct 2024 – Mar 2025 ; May 2025 – Now

Beijing Institute of Technology (BIT)

BSc in Statistics (GPA: 88/100, Major GPA: 90/100)

Beijing, China

Sept 2019 – June 2023

Relevant Coursework:

Statistics: Probability Theory, Mathematical Statistics, Multivariate Statistical Analysis, Applied Stochastic Processes.

Mathematics: Mathematical Analysis, Advanced Algebra, Optimization Methods, Function of Real Variables.

Computational Methods:

Machine Learning, Modern Natural Language Processing, Computational Social Media, Reinforcement Learning.

RESEARCH INTEREST

My long-term goal is to enable trustworthy human–AI collaboration in future agentic ecosystems. I’m particularly interested in sociotechnical challenges arising from emergent social behaviors – such as deception and collusion in multi-agent and human–AI interactions. To address these, I study both the behavioral and mechanistic aspects of LLMs to uncover failure modes and reasoning limitations. I’m especially motivated by interdisciplinary methods that distill insights from cognitive science (e.g., belief systems, decision-making) and HCI (e.g., utility functions).

PUBLICATIONS AND TECHNICAL REPORTS

- TactfulToM: Do LLMs Have the Theory of Mind Ability to Understand White Lies?  

Yiwei Liu, Emma Jane Pretty, Jiahao Huang, Saku Sugawara

In Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP 2025)

- A Visual Resolution for Syntactic Ambiguity 

Semester Project Report; conducted at IVIA lab, ETH Zürich.

- Business Model Construction Logic of Research-based Platforms (*Chapter 12*)

Co-author, contributed to data analysis; conducted under Prof. Wei Wei at Peking University.

China Machine Press, 2024 (ISBN: 978-7-111-73714-8)

RESEARCH EXPERIENCE

Interpretability Evaluation of Knowledge Editing Methods

Aug 2025 – Now

Supervisor: Dr. Thibault Laugel | AXA GO Rev Research

- Investigated the robustness and specificity of knowledge editing methods (*ROME*, *MEMIT*, *GRACE*), identifying critical limitations and under-explored unintended behaviors such as revert-back effects resembling jailbreaking and emergent hallucination patterns.
- Utilized interpretability tools (*LogitLens*, *Patchscope*) to trace how knowledge edits propagate and degrade across model layers, informing the design of more robust intervention methods for safer alignment.

Evaluating LLMs' Social Reasoning: White Lies Understanding and Theory of Mind

Oct 2024 – Mar 2025

Supervisor: Prof. Saku Sugawara

- Designed *TactfulToM*, an English ToM benchmark (100 samples, 6.7k questions) designed to evaluate LLMs' understanding of white lies through complex social scenarios via systematic templates simulating prosocial deception.
- Revealed that even state-of-the-art LLMs underperform humans in white lies understanding and reasoning, particularly in understanding the emotional motivation behind it, raising a critical alignment question: should LLMs understand white lies merely to interpret human behavior, or also to potentially generate them?

Pragmatic Communications in LLMs: From Ambiguity Detection to Clarification-Seeking

Pragmatic Reasoning in LLM Communication for Clarification-Seeking Enhancement May 2025 – Now
Supervisor: Prof. Saku Sugawara, Robin Chan

- Developing a benchmark to evaluate LLMs' clarification behavior in ambiguous dialogue contexts across domains such as medical reasoning and psychotherapy, focusing on when and how models should seek clarification instead of overconfidently responding, which often leads to incorrect task outcomes.
- Designing alignment methods grounded in pragmatic reasoning frameworks, including Rational Speech Acts for utility-driven clarification decisions and common-ground modeling for shared context tracking, to promote context-sensitive dialogue and reduce miscommunication in human–AI interaction.

Syntactic Ambiguity: Detection and Visual Resolution June 2024 – Dec 2024
Supervisors: Prof. Menna El-Assady, Prof. Frédéric Kaplan, Robin Chan

- Built a k-best constituency parser and interactive interface to detect and visualize syntactic ambiguity through alternative interpretation comparison, motivated by initial evaluations showing that LLMs struggle to reliably identify ambiguous sentences.

INTERNSHIP EXPERIENCE

Beijing Xianfeng Changqing Venture Capital (K2VC) Beijing, China
Investment Research Intern | Supervisor: Runxin Yang Dec 2022 – June 2023

- Conducted evaluations of LLM startups, tracking paradigm shifts and identifying research gaps in LLM development.

MiraclePlus (Former Y Combinator China) Beijing, China
Investment Research Intern | Supervisor: Jack Jin, Dr. Qi Lu Dec 2021 – Feb 2022

- Engaged with leading AI labs in China to track frontier developments in deep learning; performed technical evaluation on investment targets including TensorOpera AI and Thewake Systems, assessing research-to-product pathways.

JD.com Explore Academy (AI Lab) Beijing, China
Algorithm Engineer Intern | Supervisor: Dr. Chang Li June 2021 – Aug 2021

- Developed a Proof-of-Concept for JD.com × Nvidia Omniverse Virtual Retailer that collects eye-tracking and interaction data and applies multimodal deep learning methods for behavior prediction to optimize shelf layout; contributed to internal AI research strategy, focusing on multimodal and trustworthy AI for e-commerce.

SELECTED COURSE PROJECTS

Developed a Specialized LLM Tutor for EPFL Course Feb 2024 – May 2024
Modern Natural Language Processing (CS-552) | Instructor: Antoine Bosselut

- Fine-tuned Phi-2 with DoRA (PEFT) and Direct Preference Optimization (DPO) for efficient alignment; integrated Retrieval-Augmented Generation (RAG) for context-aware responses and applied quantization for deployment.

Extending Text2Image Model to Multimodal Model Sept 2023 – Dec 2023
Foundation of Digital Humanity (DH-405) | Instructor: Frédéric Kaplan

- Integrated ImageBind and Stable Diffusion into a multimodal pipeline; developed a benchmark with cross-modal inputs across major cultural groups to evaluate cultural understanding and image–text coherence, revealing limited generalization in cultural concept understanding.

HONORS AND AWARDS

- Awarded “Gintong Talent Fellowship”, Renmin University of China — 35 students selected annually 2022
- Awarded Academic Excellent Scholarship, Beijing Institute of Technology — Top 20% 2020 – 2021, 2021 – 2022

SERVICE

- Served as mentor and in the event coordination team for “Forward With Her”, a mentorship program dedicated to empowering women in STEM. Sept 2024 – Now

SKILLS

- **Programming & Frameworks:** Python, PyTorch, TensorFlow, Sklearn, Pandas, NumPy, JavaScript/HTML, Git, Linux.
- **Languages:** English (Proficient), Chinese (Native), French (Beginner), German (Beginner)